

Introduction to MapReduce

(cont.)

Rafael Ferreira da Silva

rafsilva@isi.edu

<http://rafaelsilva.com>

APP STORE

 TORNADOGUARD
From Droid Coder 2187

PLAYS A LOUD ALERT SOUND
WHEN THERE IS A TORNADO
WARNING FOR YOUR AREA.

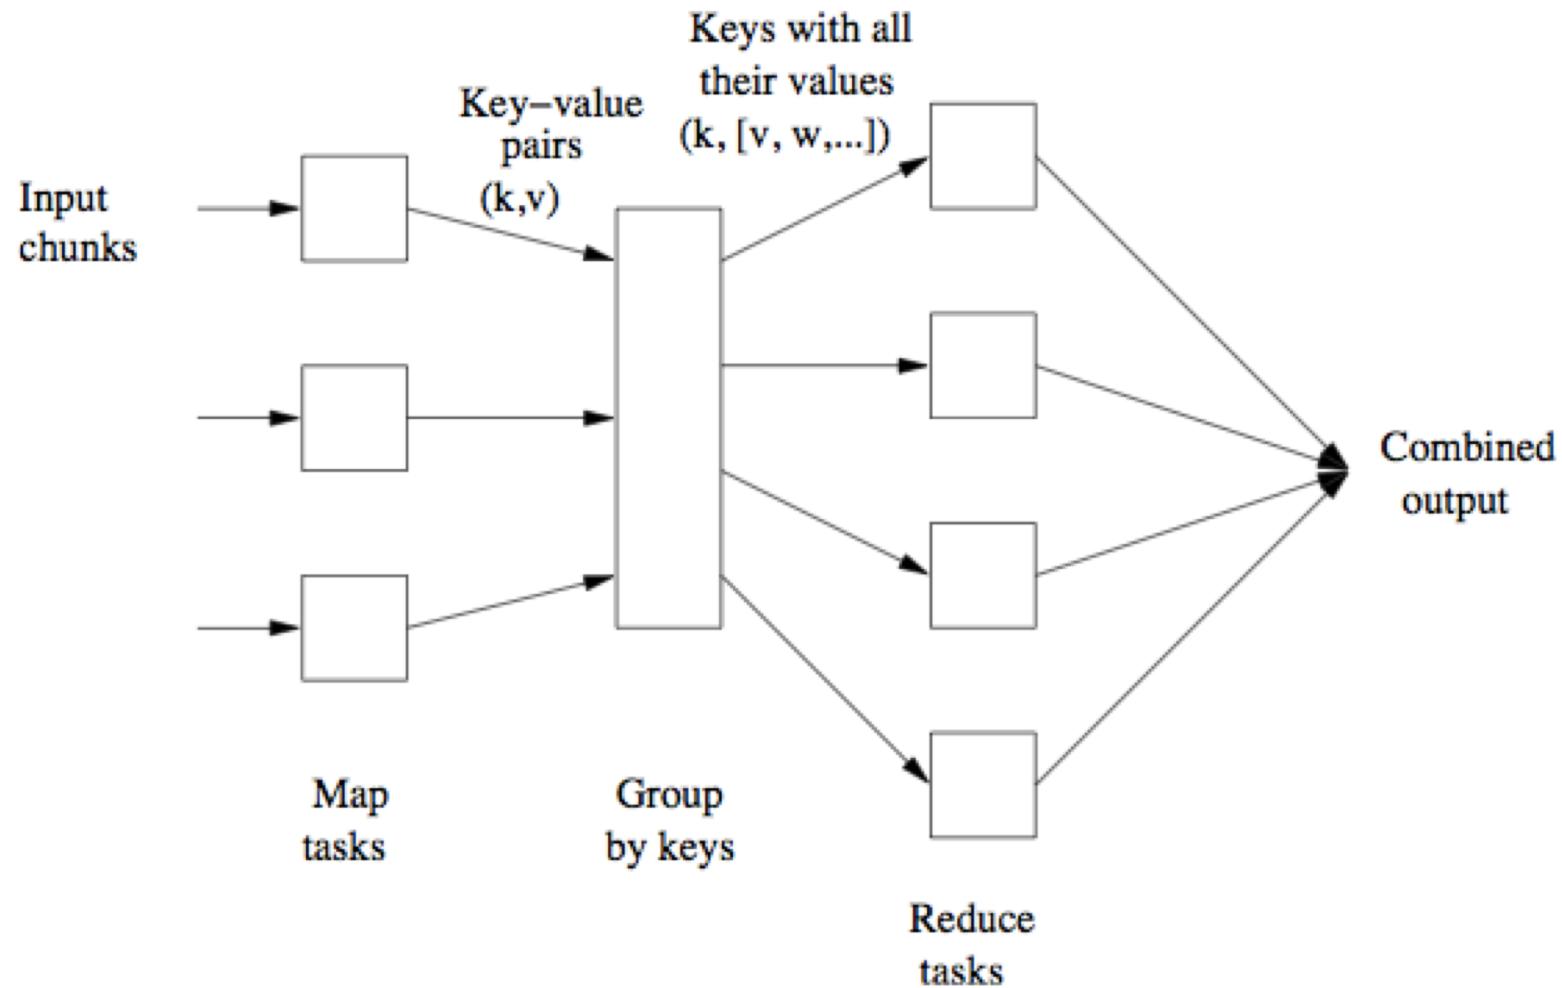
RATING: ★★★★☆
BASED ON 4 REVIEWS

USER REVIEWS:

-  ★★★★☆ GOOD UI!
MANY ALERT CHOICES.
-  ★★★★☆ RUNNING
GREAT, NO CRASHES
-  ★★★★☆ I LIKE HOW YOU
CAN SET MULTIPLE LOCATIONS
-  ★★★☆☆ APP DID NOT
WARN ME ABOUT TORNADO.

THE PROBLEM WITH AVERAGING STAR RATINGS

MapReduce: Summary



MapReduce: Examples

- Examples

```
$ docker run -it --rm -p 8888:8888 jupyter/pyspark-notebook
```

- Use file:

```
MapReduce-PySpark-Examples-2.ipynb
```

- Download **tip.json** from the yelp dataset challenge

- <https://www.yelp.com/dataset/challenge>

Combiners

- Sometimes, a Reduce function is **associative** and **commutative**
 - Can be combined in any order, with the same result
 - Can push some of what Reducers do to Map task
- Reduces network traffic for data sent to Reduce task
- Example: word count
 - Instead of producing many pairs $(w, 1), (w, 1), \dots$ can sum the n occurrences and emit (w, n)
 - Still required to do aggregation at the Reduce task for key, value pairs coming from multiple Map tasks

Example: Build an Inverted Index

Input:

tweet1, (“I love pancakes for breakfast”)
tweet2, (“I dislike pancakes”)
tweet3, (“What should I eat for
breakfast?”)
tweet4, (“I love to eat”)

Desired output:

“pancakes”, (tweet1, tweet2)
“breakfast”, (tweet1, tweet3)
“eat”, (tweet3, tweet4)
“love”, (tweet1, tweet4)
...

Map task:

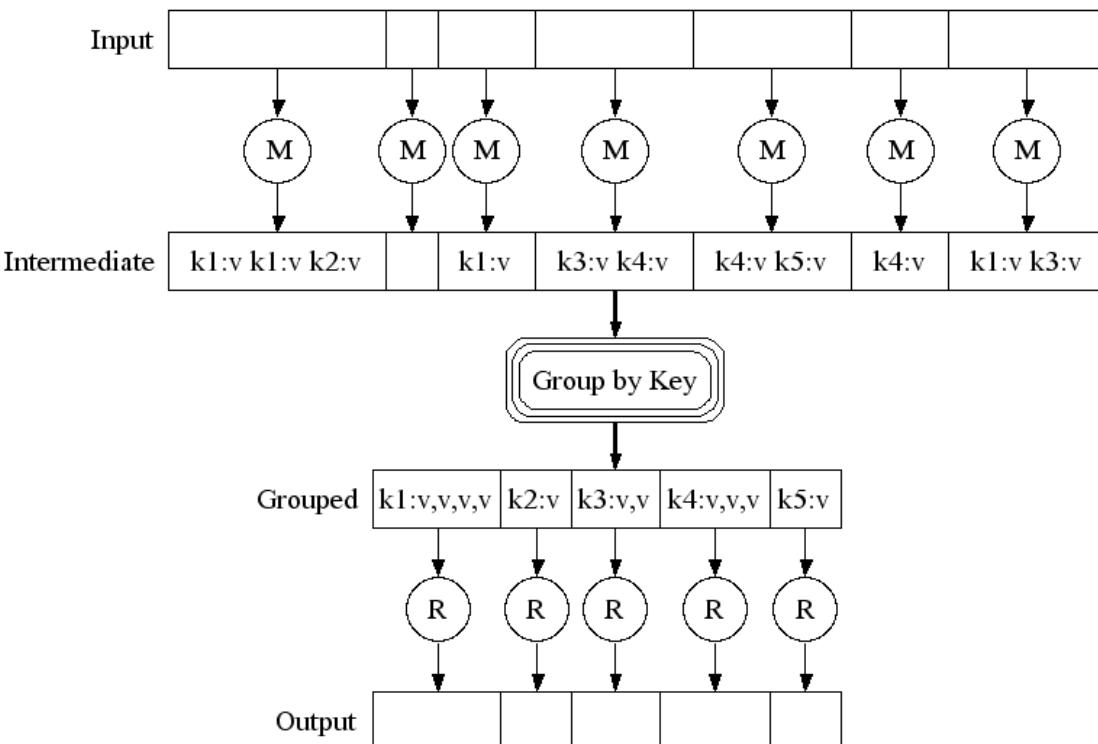
For each word, emit (word, tweetID) as intermediate key-value pair

Reduce task:

Reduce function then just emits key and list of tweetIDs associated with that key

MapReduce Environment

- MapReduce environment takes care of:
 - **Partitioning** the input data
 - **Scheduling** the program's execution across a set of machines
 - Performing the **group by key** step
 - In practice this is the bottleneck
 - Handling machine **failures**
 - Managing required inter-machine **communication**



Data Flow

- **Input and final output** are stored on a **distributed file system (HDFS)**
 - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- **Intermediate results** are stored on **local FS** of Map and Reduce workers
- Output is often **input** to another MapReduce task

Coordination: Master

- Master node takes care of coordination
 - Task status: (idle, in-progress, completed)
 - Idle tasks get scheduled as workers become available
 - When a map task completes, it sends the master the location and sizes of its intermediate files
 - one for each reducer
 - Master pushes this info to reducers
- Master pings workers periodically to detect failures

Dealing with failures

- **Map worker failure**
 - Map tasks completed or in-progress at worker are reset to idle and rescheduled
 - Reduce workers are notified when task is rescheduled on another worker
- **Reduce worker failure**
 - Only in-progress tasks are reset to idle, why?
 - Reduce task is restarted
- **Master failure**
 - MapReduce task is aborted and client is notified

How many Map and Reduce jobs?

- M map tasks, R reduce tasks
- **Rule of a thumb**
 - Make M much larger than the number of nodes in the cluster
 - One DFS chunk per map task is common
 - Improves dynamic load balancing and speeds up recovery from worker failures
- **Usually R is smaller than M**
 - Because output is spread across R files
- **Google example**
 - Often use 200,000 map tasks, 5000 reduce tasks on 2000 machines

Task granularity and pipelining

- **Fine granularity tasks:** map tasks >> machines
 - Minimizes time for fault recovery
 - Can do pipeline shuffling with map execution
 - Multiple mapper on same machine. One working, the other output/shuffling
 - Better dynamic load balancing

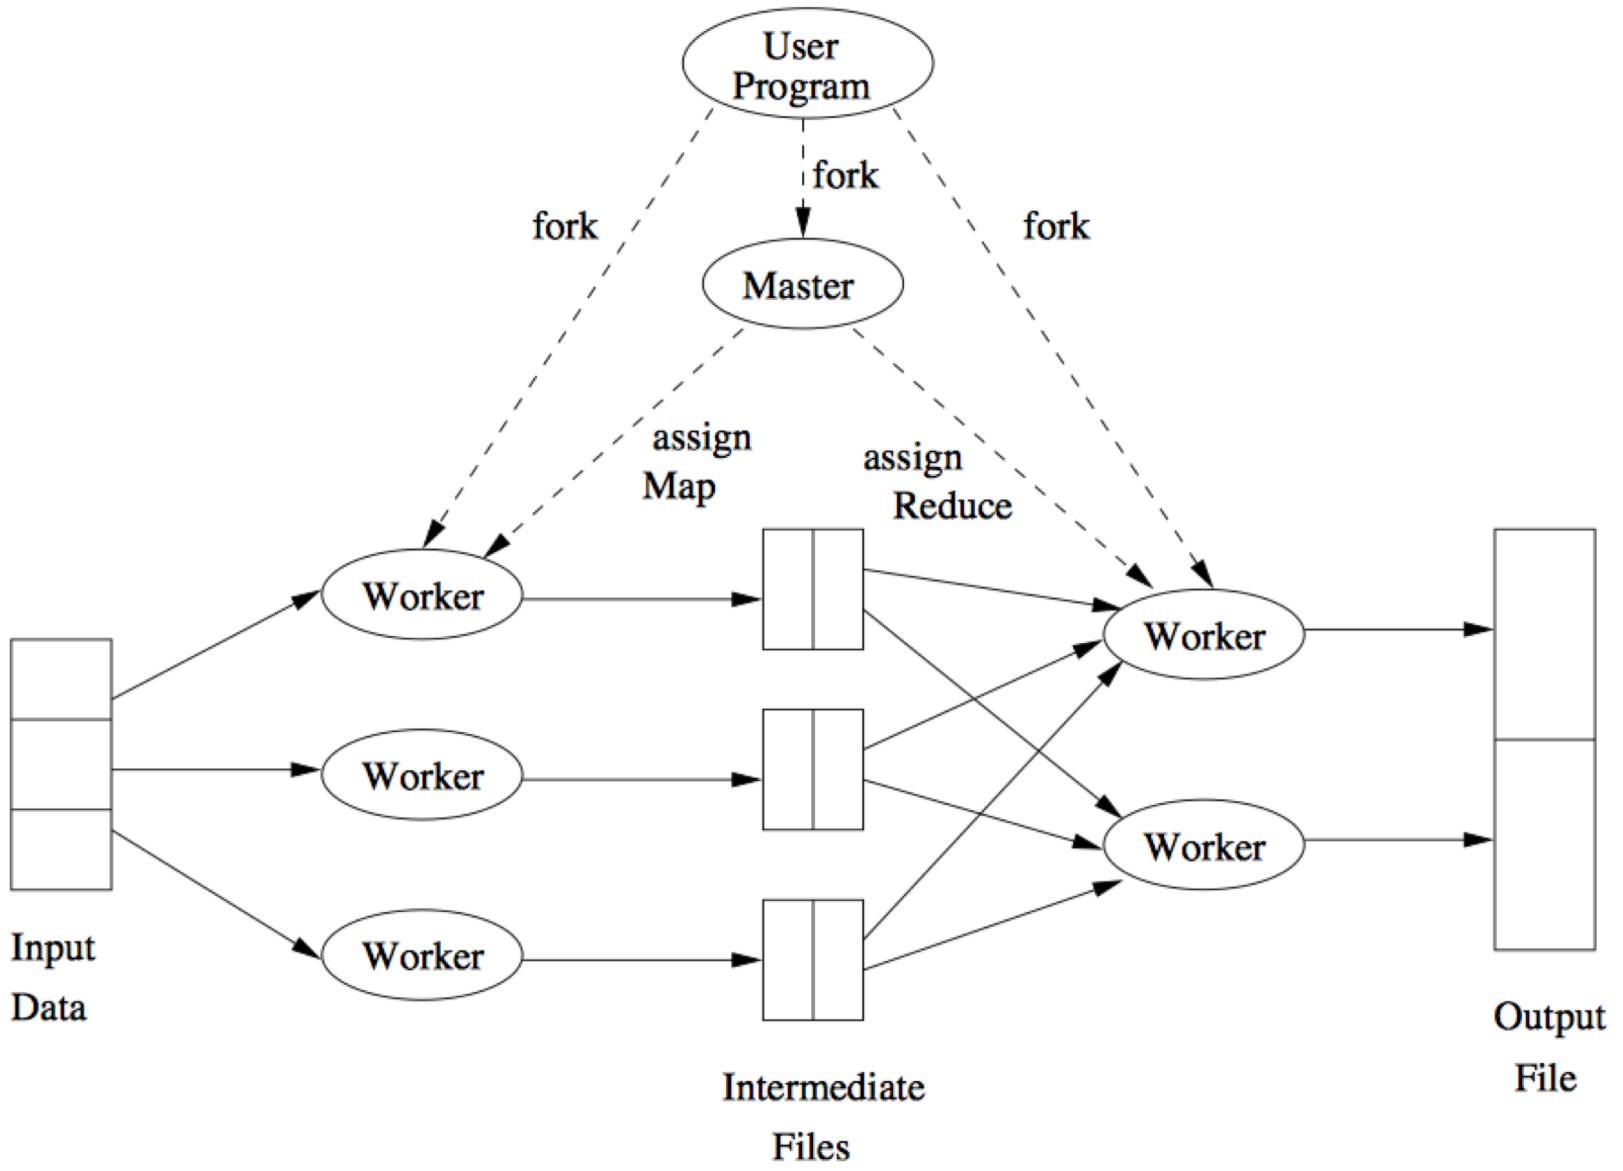


Figure 2.3: Overview of the execution of a MapReduce program

Refinements: Backup Tasks

- **Problem**

- Slow workers significantly lengthen the job completion time
 - Other jobs on the machine
 - Bad disks
 - Weird things

- **Solution**

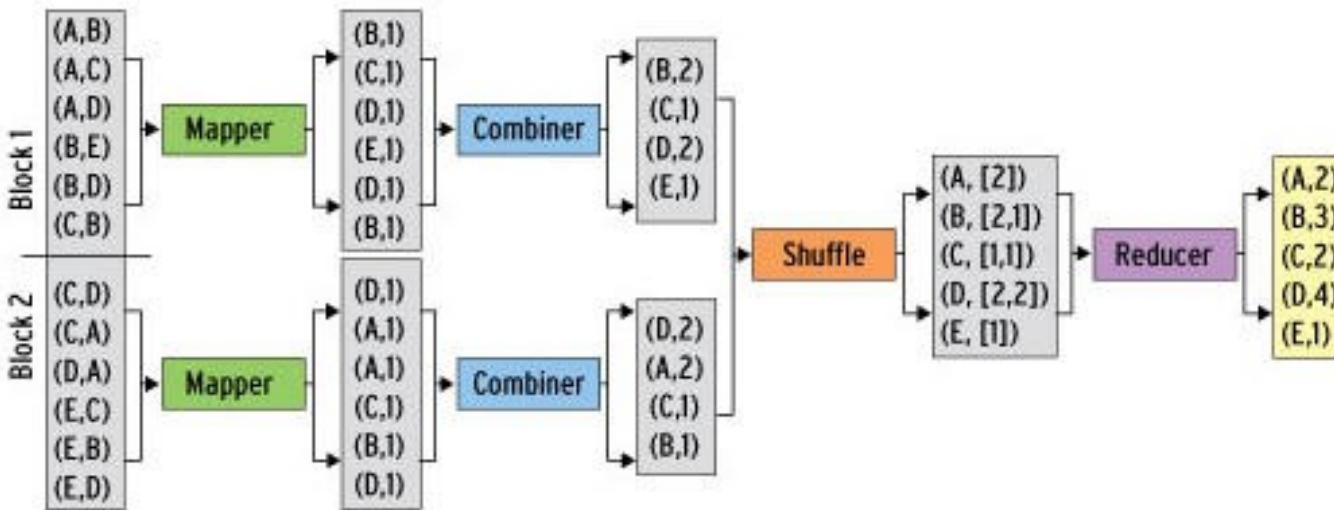
- Near end of phase, spawn backup copies of tasks
 - Whichever one finishes first “wins”

- **Effect**

- Dramatically shortens job completion time

Refinements: Combiners

- Back to our word counting example
 - **Combiner** combines the values of all keys of a single mapper ...



- Much less data needs to be copied and shuffled!
- Works if reduce function is **commutative and associative**

Refinement: Partition function

- Want to control how keys get partitioned
 - Inputs to map tasks are created by **contiguous splits of input file**
 - Reduce needs to ensure that **records with the same intermediate key** end up at the same worker
- **System uses a default partition function**
 - **hash(key) mod R**
- Sometimes useful to override the hash function
 - E.g., want to have alphabetical or numeric ranges going to different Reduce tasks (sorting)
 - E.g., **hash(hostname(URL)) mod R** ensures URLs from a host end up in the same output file

Problems Suited for MapReduce

General Characteristics of Good Problems for MapReduce

- Data set is truly “big”
 - **Terabytes**, not tens of gigabytes
 - Hadoop/MapReduce designed for terabyte/petabyte scale computation
- Most real-world problems process less than 100 Gbytes of input
 - Microsoft, Yahoo: median job under 14 GB
 - Facebook: 90% of jobs under 100 GB

Source: “To Hadoop or Not to Hadoop” by Anand Krishnaswamy 8/13/2013

General Characteristics of Good Problems for MapReduce

- Don't need **fast response time**
- When submitting jobs, Hadoop latency can be **a minute**
- Not well-suited for problems that require **faster response time**
 - online purchases, transaction processing
- A good **pre-computation engine**
 - E.g., pre-compute related items for every item in inventory

Source: "To Hadoop or Not to Hadoop" by Anand Krishnaswamy 8/13/2013

General Characteristics of Good Problems for MapReduce

- Good for applications that work in **batch mode**
- Jobs run without human interaction, intervention
- **Runs over entire data set**
 - Takes time to initiate, run; shuffle step can be time-consuming
- Does not support real time applications well: sensor data, real-time advertisements, etc.
- Does not provide good support for **random access to datasets**
 - Extensions: Hive, Dremel, Shark, Amplab

Source: "To Hadoop or Not to Hadoop" by Anand Krishnaswamy 8/13/2013

General Characteristics of Good Problems for MapReduce

- Best suited for data that can be expressed as **key-value pairs** without losing context, dependencies
- **Graph data harder to process using MapReduce**
 - Implicit relationships: edges, sub-trees, child/parent relationships, weights, etc.
- Graph algorithms may need information about the entire graph for each iteration
 - Hard to break into independent chunks for Map tasks
- Alternatives: Google's Pregel, Apache Giraph

Source: "To Hadoop or Not to Hadoop" by Anand Krishnaswamy 8/13/2013

General Characteristics of Good Problems for MapReduce

Other problems/data ***Not*** suited for MapReduce

- Tasks that need results of intermediate steps to compute results of current step
 - **Interdependencies among tasks**
 - Map tasks must be independent
- Some machine learning algorithms
 - Gradient-based learning, expectation maximization

Source: “To Hadoop or Not to Hadoop” by Anand Krishnaswamy 8/13/2013

General Characteristics of Good Problems for MapReduce

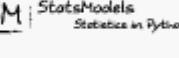
Summary: Good candidates for MapReduce

- Jobs that have to process **huge quantities of data** and either **summarize or transform the contents**
- Collected data has elements that can easily be captured with an identifier (key) and corresponding value

Source: "To Hadoop or Not to Hadoop" by Anand Krishnaswamy 8/13/2013

Useful tools for Data Science

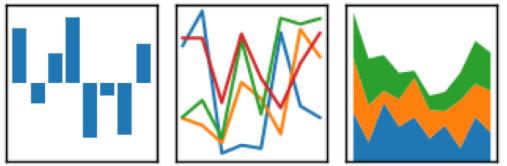
Python libraries for data science

Library Name	Type	Library Name	Type	Library Name	Type
 matplotlib	Visualization	 eli5	Machine learning	 Keras	Deep learning
 Bokeh	Visualization	 SciPy	Data wrangling	 dist-keras	Distributed deep learning
 plotly	Visualization	 NumPy	Data wrangling	 elephas	
 Seaborn	Visualization	 pandas	Data wrangling	 spark-deep-learning	
 pydot	Visualization	 StatsModels	Statistics	 Natural Language ToolKit	NLP
 scikit-learn	Machine learning	 TensorFlow	Deep learning	 spaCy	NLP
XGBoost LightGBM CatBoost	Machine learning	 PYTORCH	Deep learning	 gensim	NLP
				 Scrapy	Data scraping

<https://activewizards.com/blog/top-20-python-libraries-for-data-science-in-2018/>

Pandas

pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



<https://pandas.pydata.org>

- Python library that provides high-level data structures and a vast variety of tools for analysis
- Ability to translate rather **complex operations** with data into one or two commands
- Pandas contains many **built-in methods**
 - Grouping
 - Filtering
 - Combining data
 - Missing data
 - Time-series
 - Hierarchical indexing

Scikit-learn

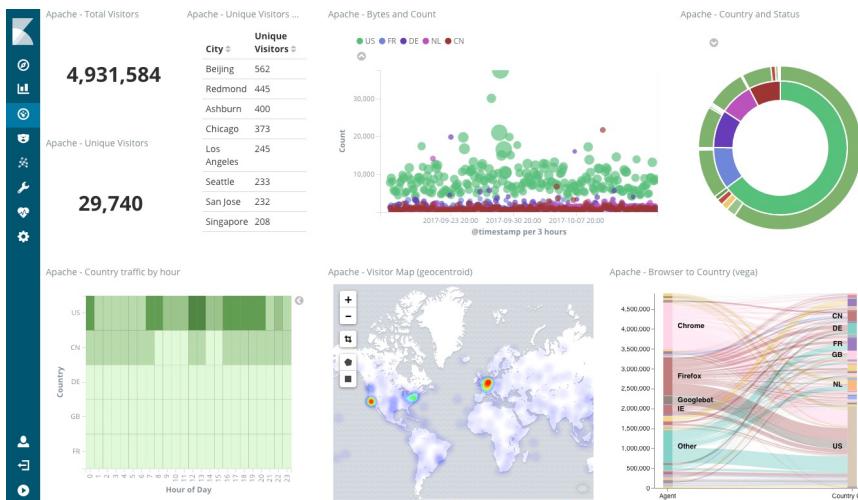
- Python module based on **NumPy** and **SciPy**
- Provides algorithms for many standard **machine learning** and **data mining** tasks
 - Classification
 - Regression
 - Clustering
 - Dimensionality reduction
 - Model selection
 - Preprocessing

Elasticsearch



<https://www.elastic.co/products/elasticsearch>

- Open-source, RESTful, distributed search and analytics engine
 - **Real-time** data and real-time analytics
 - Scalable, high-availability, multi-tenant
 - **Full text** search



<https://www.elastic.co/products/kibana>



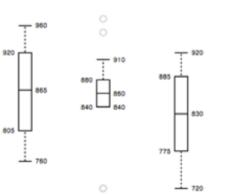
<https://grafana.com>

D3.js

- Data Driven Documents
 - JavaScript library for manipulating documents based on data

<https://d3js.org>

Box Plots



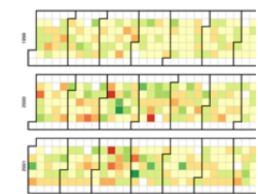
Bubble Chart



Bullet Charts



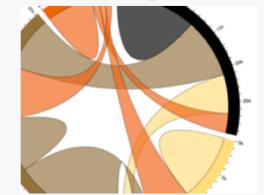
Calendar View



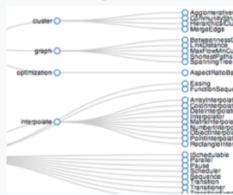
Non-contiguous
Cartogram



Chord Diagram



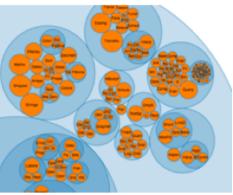
Dendrogram



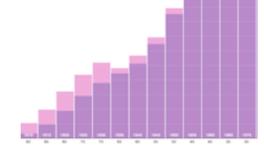
Force-Directed
Graph



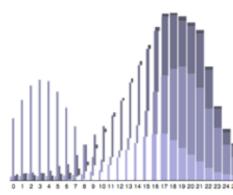
Circle Packing



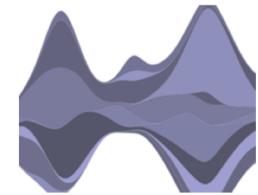
Population Pyramid
2000



Stacked Bars



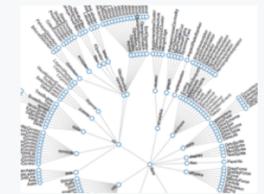
Streamgraph



Sunburst



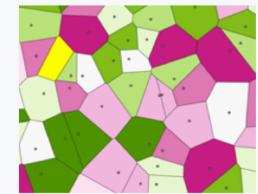
Node-Link Tree



Treemap



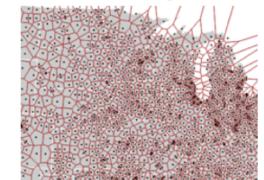
Voronoi Diagram



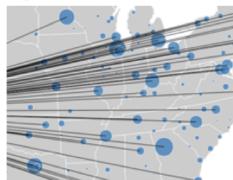
Hierarchical Edge
Bundling



Voronoi Diagram



Symbol Map



Parallel Coordinates

